



arm

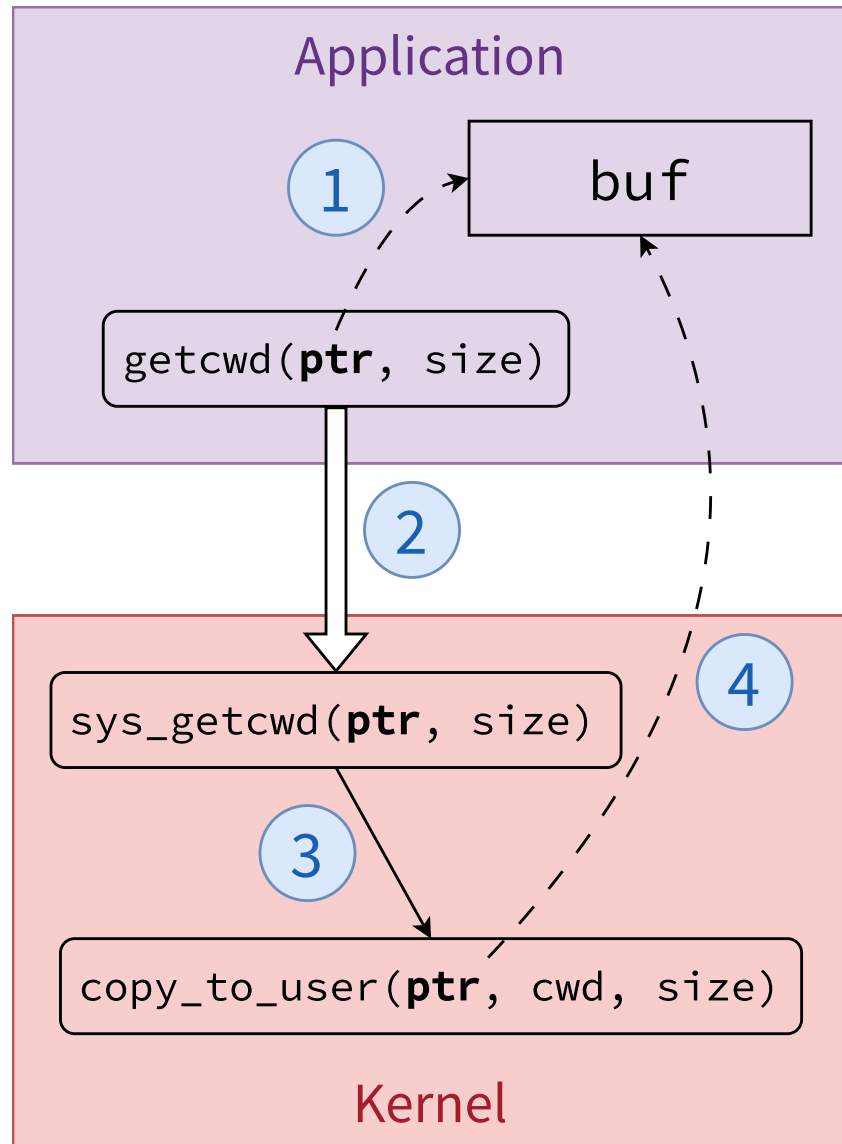
Morello Linux kernel overview

Kevin Brodsky

The pure-capability userspace ABI (PCuABI)

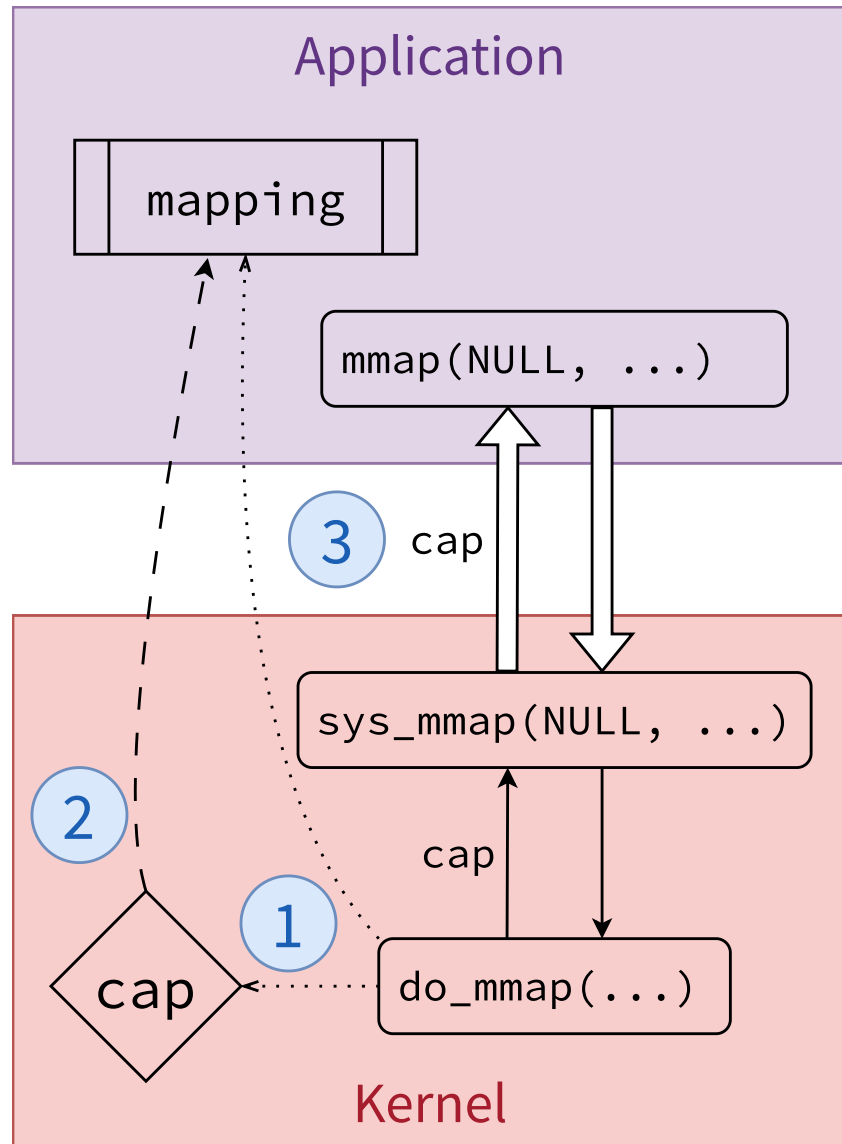
- New kernel-user ABI required
 - **Functional** angle: userspace uses larger, "special" pointers
 - **Security** angle: enforcement of capability properties
-
- Input pointers (user → kernel)
 - Most common (syscall arguments)
 - Output pointers (kernel → user)
 - Mostly address space management (`mmap()`, ...)

Kernel-user interactions: input pointers



1. `ptr` is a capability allowing access to `buf` only
2. PCuABI syscall: `ptr` passed as a capability (`c0`)
3. `ptr` is **propagated** as a capability
4. Capability-based uaccess: `buf` is accessed via the `ptr` capability
 - Exception triggered if `ptr` does not authorise the access (`-EFAULT`)

Kernel-user interactions: output pointers



1. New mapping and capability created
2. `cap` grants access to `mapping`
 - Minimal bounds and permissions
 - `cap` **owns** `mapping` (VMem permission)
→ allows calling `mprotect()`, etc.
3. `cap` is returned to userspace

Capability propagation: hybrid approach

- All **user pointers** become capabilities through annotations
 - C extension: `void * __capability`
 - Leveraging `__user` — some fixups needed as `__user` prefixes `*`
- Primitive types (in-kernel ABI) **unchanged**
 - Kernel pointers, `long` still 64-bit
- Strict separation between kernel and user pointers
 - New APIs to manipulate user pointers: `<linux/user_ptr.h>`
 - **Address \neq pointer**

Representing user pointers

- User pointers often represented as integers

```
long (*unlocked_ioctl) (struct file *, unsigned int, unsigned long);
```

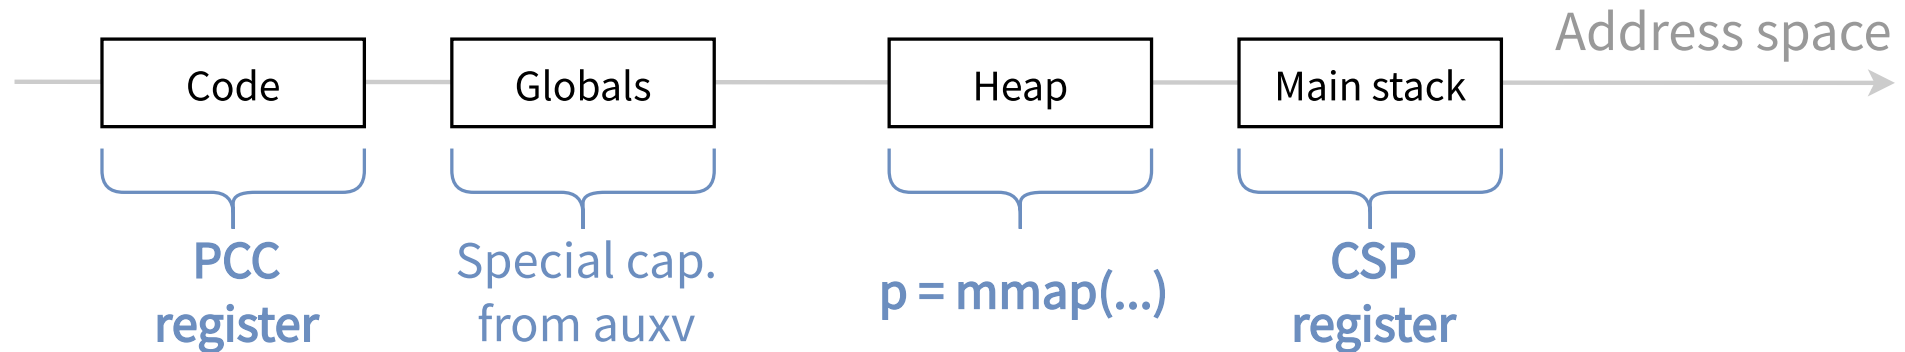
```
struct io_uring_sqe {  
    ...  
    __u64 addr; /* pointer to buffer or iovecs */  
    ...  
};
```

- Not capability-friendly: `(long)uptr` truncates `uptr`
- New types required — ABI-dependent definition

`unsigned long` → `user_uintptr_t`

`__u64` → `__kernel_uintptr_t`

Providing root capabilities



- Mainly: `mmap()`, `mremap()`
 - Return a capability with minimal bounds and permissions
 - Special handling of input pointers
- Initial process environment
 - PCC (code capability), CSP (capability stack pointer)
 - Arguments (`argv`) and env. variables (`envp`): arrays of capabilities
 - New `auxv` entries for special root capabilities (mainly for relocation processing)



compat64: 64-bit compat layer

- Standard 64-bit ABI provided via compat
 - Instead of 32-bit (not present on Morello HW)
- compat → native pointer conversion required
 - 64-bit address → valid capability
 - User memory accessed via a capability in any case
 - **Strict usage** of `compat_ptr()`
 - Pointer arguments converted directly in syscall wrapper

Application

```
char buf[128];
struct iovec buf_iov = {
    .iov_base = &buf,
    .iov_len = 128
};
readv(fd, &buf_iov, 1);
```

Kernel

```
sys_readv(int fd, struct iovec __user *iov,
          int iovcnt);
```

- `iov` capability created by the syscall wrapper
- `sys_readv()` uses `compat_ptr(iov[0].iov_base)`

compat64: pain points

- Widespread assumption that compat is 32-bit
 - We want most of the existing compat code... but not all of it (e.g. 32-bit time types)
 - `unsigned int` → `compat_ulong_t`
 - Sometimes: just use native handler
- Additional handling where types have been enlarged for native pointers
 - **Any uapi struct change requires adding compat handling**
 - Typically: `__u64` → `__kernel_uintptr_t`
 - Layout conversion not always doable upfront (e.g. `union bpf_attr`)
 - compat ioctl handlers must always convert input pointers

Morello Linux kernel

- Morello Linux kernel fork [hosted on morello-project.org](https://morello-project.org)
 - Mainline-based (currently 6.7) + ~500 patches
 - Support for Morello and pure-capability userspace (**hybrid** approach)
 - Selection of drivers available in PCuABI
 - 64-bit compat → **major effort**
- Pure-capability kernel-user ABI specification — [wiki page on morello-project.org](https://morello-project.org/wiki)
 - Refined over 3+ years, mostly stable
 - Extensive, many subtleties around `mmap()` and address space management in general

arm

Thank you
Danke
Merci
ありがとう
谢谢
Gracias
Tack
Takk
Kiitos
감사합니다
धन्यवाद
تشکر



The Arm trademarks featured in this presentation are registered trademarks or trademarks of Arm Limited (or its subsidiaries) in the US and/or elsewhere. All rights reserved. All other marks featured may be trademarks of their respective owners.

www.arm.com/company/policies/trademarks

